



منة الله اشرف على محمد 2205252

Report:

1. Objective of the Code and the Experiment

In this code, I build a simple experiment using a **Graph Neural Network (GraphSAGE)** to do **node classification** on a small graph.

The main idea is that I have a group of users in a network (for example, a social network). Some of them are **benign (normal/safe users)** and some are **malicious (harmful users)**, and I want to train a model that can predict whether each user is benign or malicious based on:

- The **features** of each user
- The **connections (edges)** between users in the graph

I use **PyTorch Geometric** to take advantage of the ready-made **GraphSAGE layers**.

2. Preparing the Graph Data

2.1 Node Features x

First, I define the tensor x, which represents the **features of each node** in the graph:

```
x = torch.tensor(  
    [  
        [1.0, 0.0], # Node 0 (benign)  
        [1.0, 0.0], # Node 1 (benign)  
        [1.0, 0.0], # Node 2 (benign)  
        [0.0, 1.0], # Node 3 (malicious)  
        [0.0, 1.0], # Node 4 (malicious)  
    ]
```



```
[0.0, 1.0] # Node 5 (malicious)  
],  
dtype=torch.float,  
)
```

- I have **6 nodes** (from 0 to 5).
- Each node has **2 features**:
 - [1, 0] = benign user
 - [0, 1] = malicious user

So here I'm giving the model a simple **one-hot encoding** that tells it there are only two types of nodes, benign and malicious, and each type has its own fixed feature vector.

2.2 Defining the Edges edge_index

Then I define edge_index, which represents the **connections (edges)** between the nodes:

```
edge_index = (  
    torch.tensor(  
        [  
            [0, 1],  
            [1, 0],  
            [1, 2],  
            [2, 1],  
            [0, 2],  
            [2, 0],  
        ]  
    )
```



```
[3, 4],  
[4, 3],  
[4, 5],  
[5, 4],  
[3, 5],  
[5, 3],  
[2, 3],  
[3, 2],  
],  
dtype=torch.long,  
)  
.t()  
.contiguous()  

```

I am treating the graph as **undirected**, so for every edge (u, v) I also add (v, u) :

- Nodes **0, 1, 2** form a **complete triangle** → this is the **benign group** that is well connected internally.
- Nodes **3, 4, 5** form another triangle → this is the **malicious group**.
- Finally, I add **one edge between 2 and 3** to create a **connection between the benign world and the malicious world** (a benign user connected to a malicious user).

This makes the scenario a bit more interesting, because the model sees that **Node 2 (benign)** is connected to nodes in the malicious cluster.



2.3 Defining the Labels y

Next, I define the labels (the **true class** for each node):

```
y = torch.tensor([0, 0, 0, 1, 1, 1], dtype=torch.long)
```

- 0 = benign
- 1 = malicious

So:

- Nodes 0, 1, 2 → benign → label = 0
- Nodes 3, 4, 5 → malicious → label = 1

2.4 Packing Everything into a Data Object

PyTorch Geometric expects the data to be put into a Data object:

```
data = Data(x=x, edge_index=edge_index, y=y)
```

Here:

- data.x = node features
- data.edge_index = edges
- data.y = ground-truth labels

This data object is what I pass to the model during **training** and **evaluation**.

3. Designing the GraphSAGE Model

3.1 Defining the GraphSAGENet Class

I build a simple model with **two GraphSAGE layers**:

```
class GraphSAGENet(torch.nn.Module):
```

```
    def __init__(self, in_channels, hidden_channels, out_channels):
```



```
super(GraphSAGENet, self).__init__()  
  
self.conv1 = SAGEConv(in_channels, hidden_channels)  
  
self.conv2 = SAGEConv(hidden_channels, out_channels)  
  
  
def forward(self, x, edge_index):  
  
    x = self.conv1(x, edge_index)  
  
    x = F.relu(x)  
  
    x = self.conv2(x, edge_index)  
  
    return F.log_softmax(x, dim=1)
```

Architecture explanation:

- `in_channels = 2` → because each node has 2 features.
- `hidden_channels = 4` → I chose a 4-dimensional hidden embedding.
- `out_channels = 2` → I have 2 classes (benign and malicious), so the model outputs 2 scores (logits) per node.

3.2 What Happens Inside forward (Step by Step)

1. `self.conv1(x, edge_index)`
 - GraphSAGE takes the node features and:
 - Reads each node's own features.
 - Aggregates information from its neighbors based on `edge_index`.
 - Produces a new **embedding** for each node that combines its own features + the neighbors' features.
2. `F.relu(x)`



- I apply the **ReLU activation function** to add non-linearity and improve the expressive power of the model.
3. `self.conv2(x, edge_index)`
- I apply a second GraphSAGE layer on top of the first-layer embeddings.
 - The output now has shape [num_nodes, out_channels], i.e., for each node I get 2 values (one for each class).
4. `F.log_softmax(x, dim=1)`
- I convert the logits into **log-probabilities** over the two classes.
 - I use `log_softmax` because later, in the loss function, I use `F.nll_loss`, which expects log-probabilities.

4. Training Phase (Training Loop)

4.1 Creating the Model and the Optimizer

```
model = GraphSAGENet(in_channels=2, hidden_channels=4, out_channels=2)  
optimizer = torch.optim.Adam(model.parameters(), lr=0.01)
```

- I choose the **Adam** optimizer because it usually works well out of the box without much tuning.
- The learning rate `lr = 0.01` is a reasonable value for such a small and simple experiment.

4.2 The Training Loop

```
model.train()  
  
for epoch in range(50):  
    optimizer.zero_grad()  
    out = model(data.x, data.edge_index)
```



```
loss = F.nll_loss(out, data.y)  
loss.backward()  
optimizer.step()
```

At each epoch, the following steps happen:

1. `optimizer.zero_grad()`
 - o I reset the gradients from the previous iteration.
2. `out = model(data.x, data.edge_index)`
 - o I run the **forward pass**:
 - The model receives the features and the edges.
 - It outputs `out`, which are the log-probabilities for each node over the 2 classes.
3. `loss = F.nll_loss(out, data.y)`
 - o I compute the **loss** between the model predictions and the true labels `y`.
 - o Since the model outputs log-softmax, `nll_loss` is the correct choice here.
4. `loss.backward()`
 - o I perform **backpropagation**, computing gradients for all model parameters.
5. `optimizer.step()`
 - o I update the model weights using the gradients and the learning rate.

I train the model for **50 epochs**, which is more than enough for such a tiny graph with only 6 nodes.



5. Evaluation and Reading the Results

After training, I switch the model to evaluation mode:

```
model.eval()  
  
pred = model(data.x, data.edge_index).argmax(dim=1)  
  
print("Predicted labels:", pred.tolist())
```

- `model.eval()`
 - This turns off training-specific behaviors (like Dropout or BatchNorm if they existed), so evaluation is stable.
- `model(data.x, data.edge_index)`
 - I run a forward pass again to get the log-probabilities for each node.
- `.argmax(dim=1)`
 - I take the class with the highest log-probability for each node:
 - 0 → predicted as benign
 - 1 → predicted as malicious

If training goes well, I expect to see something like:

Predicted labels: [0, 0, 0, 1, 1, 1]

This means the model has successfully learned to **separate the benign group from the malicious group** and classify all 6 nodes correctly based on both their features and their connections in the graph.